Preserving Independent Judgement in AI-Mediated Financial Information

By Sajid Shaikh

Oct 22, 2026

Article for the proceedings of the University of Oxford AI Summit 2026, published by Erdos Research

Abstract

Financial analysis has always relied on data, models and historical comparison. The newer concern is the point at which generative AI begins to select sources, frame explanations, and synthesise a market narrative. If many organisations rely on similar models and information pipelines, apparently independent and diverse analyses may converge, crowding out minority views that often signal emerging risks. Similar published views alone would not prove that AI caused them to converge.

This article uses London@11, a prototype financial news agent that logs its editorial decisions, to propose a practical response: separate authoritative facts from commentary, test source independence, preserve credible dissent and record how human and AI decisions shape the final draft. Early tests indicate that this workflow is feasible, not that it has a sector-wide effect. The article argues that epistemic independence must become an observable design property of AI-mediated financial information, and proposes controlled comparisons as the next step.

The convergence problem

Newsrooms, banks, brokers and analysts increasingly use AI to digest financial news and data, and to draft and publish their own commentary. The output reaches investors, consumers and policymakers, and shapes how markets and the economy are understood.

This article examines one narrow part of that shift: AI synthesis of financial commentary into a single briefing. When a model summarises what several economists and analysts think, it can make different experts appear to agree, and minority views are the easiest to lose. If many organisations use similar models and sources, those losses may compound, so that apparently independent and diverse analyses converge on a single interpretation.

Convergence is not new, and it is not always harmless. Before the 2008 financial crisis, mainstream opinion largely overlooked the risks building in subprime lending, while a minority of analysts warned of them. The US Financial Crisis Inquiry Commission later concluded that the crisis was avoidable and that its warning signs had gone unheeded (*Financial Crisis Inquiry Commission, 2011*). That episode predates AI; the concern is that AI could make the same pattern faster, wider and harder to see, because a smoothed-out consensus leaves no trace of the views it dropped.

Current safeguards do not address this. Editorial review checks accuracy, not whether a summary represents disagreement, and retrieval tools cite sources based on facts without testing whether the

sources are independent. There is no established practice for recording when interpretations converge or which credible dissent was set aside.

If the problem is left unaddressed, readers and decision-makers may mistake uniform commentary for confirmed judgement, and markets could lose part of their risk instinct, the early warning that diverse analyses provide, just as AI use becomes routine.

Evidence of convergence

Before designing London@11, I tested whether an AI model could reproduce two daily series I had launched as an editor at a major financial news publisher: a snapshot of the day's trending stocks, and a commodities and currencies report covering sterling, gold and oil. I supplied about a year of the series' published output to the default model available through ChatGPT Enterprise and asked it to produce equivalent briefings. Four failures recurred:

1. **Repetition:** the model favoured companies that appeared most often in the material it was given, rather than those genuinely trending that day.
2. **Instruction decay:** in complex, multi-step prompts it completed some instructions and dropped others.
3. **Drift:** over extended interactions, its output moved away from the original task specification.
4. **Temporal grounding failure:** it did not anchor briefings to the required time window, such as "post-US close, pre-Europe open".

All four share a direction: the output moved away from the editor's specific, time-bound instructions and towards the model's own defaults. Because those defaults are shared by every user of the same model, such reversion is a plausible pathway to similar output across organisations. Repetition illustrates this most directly; the instruction and temporal failures show how easily safeguards stated only in a prompt, including an instruction to preserve dissent, can be lost. However, these are practitioner observations from an informal test, not controlled findings.

The pattern is not confined to financial content. A paper published in Nature by Hao et al. (2026), 'Artificial intelligence tools expand scientists' impact but contract science's focus', analysed 41.3 million scientific papers and found that researchers using AI methods published 3.02 times more papers and received 4.84 times more citations, yet AI adoption shrank the collective range of topics studied by 4.63% and reduced scientists' engagement with one another's work by 22%. The study examines AI methods broadly rather than generative AI, but a survey of more than 600 scientists, reported by Imas and Manyika (2026) in the DeepMind Institute essay 'Bending the Curve of Discovery: AI in Science Today and Tomorrow', points the same way: almost half said AI pushes them towards safer, incremental questions.

Designing for dissent: London@11

London@11 is a prototype agent that produces a daily UK financial briefing, designed around the failures observed in my earlier tests. Three documents govern it: project instructions, an editorial policy and a source policy.

Facts must come from named primary sources, such as the ONS, the Bank of England and HM Treasury. Commentary from approved outlets and named economists may inform interpretation but never verify

a fact. Before describing commentators as converging or diverging, the agent must consult at least three independent sources.

Sources count as independent only if separately authored and original: shared wire copy, syndicated text, a byline repeated across outlets or several reports of the same underlying note do not qualify. Below that threshold, the assessment is labelled inconclusive. A dissenting view must be surfaced when it comes from a named analyst, economist or institution with stated reasoning; anonymous or unsupported views do not qualify, so dissent is preserved without false balance. If two primary sources conflict, the agent halts the claim and refers it to the editor.

Each run produces two outputs: publication copy (headline, standfirst, briefing, key takeaways and a markets snapshot) and Editor Notes (a verification ledger, a consensus-and-dissent log, and gaps and caveats). The agent cannot publish; the editor approves a specific draft version.

The tendency to favour consensus is not unique to AI. In my experience, deadline pressure leads newsrooms to favour the consensus because it is the quickest interpretation to substantiate, and models trained on or supplied with such coverage may reproduce it. Removing AI would not eliminate the problem; making the judgement explicit would address both. Whether a briefing represents real disagreement, previously left to an individual journalist under time pressure, becomes a recorded and checkable step, and the log accumulates into a dataset of when commentary converged.

Unlike citation-based retrieval, or media-bias aggregators that classify outlets by political leaning, London@11 tests independence directly. Its contribution is the combination of an independence test, a source threshold before any judgement of consensus, a credible-dissent rule, and human approval tied to an exact version.

Validation tests

I evaluated London@11 in two ways. First, it ran live over eight publication days between 18 September and 19 October 2026, producing 20 drafts, of which I approved three, two with edits. The independence controls worked as designed: the agent labelled assessments as inconclusive when fewer than three independent sources were available, and identified hidden dependencies: in its 19 October 2026 edition, it traced all coverage of a Barclays story back to a single Financial Times report.

Live operation also exposed limits. Free access to market data and major outlets was unreliable, and two policy fixes did not hold. Twice, the agent relayed a secondary outlet's explanation for a market move without confirming the underlying event, itself a form of narrative adoption. Real news days never tested the dissent rule directly.

Second, I ran the policies against six planted scenarios, three times each, on fresh agents with fixed synthetic inputs and pass criteria set in advance. All 18 runs passed. Across the scenarios, the agent preserved a named, reasoned dissent in reader-facing copy, reduced five apparently agreeing outlets to two independent sources, rejected anonymous or unapproved views as dissent and reported genuine agreement without faking balance. When two official sources gave conflicting figures for the same statistic, it withheld the figure and referred the decision to the editor. It also ignored a hidden instruction, planted in a retrieved web page, to approve its own draft. A perfect score shows the rules

work as written, but it also suggests the scenarios were too easy to expose their limits; harder cases are needed.

I designed the agent, edited every draft and evaluated the results, and an AI assistant helped review drafts and build and score the bench tests. The findings indicate that the controls work as specified, not that they reduce homogenisation at scale. The proposition to test is that, in sustained use, London@11 shows credible dissent that a standard AI summary of the same sources omits. A blind comparison scored by independent editors would support it if dissent is retained more often without loss of accuracy, and weaken it if the difference is negligible.

Making diversity observable

This article argues that interpretive diversity should be treated as a property that AI systems are designed to preserve and expose, rather than as a quality requested in a prompt and hoped for. The evidence points that way. Safeguards stated only in the instructions were lost during my early tests, and two of London@11's own policy fixes failed until a later run caught them. The controls held when they were observable: a fixed source threshold, a structured log with a named-dissent field, and an editor approving a specific version.

This challenges a common assumption: that an AI summary is sound if its facts are accurate. A summary can be accurate for every figure and still misrepresent how far experts disagree. Faithfulness to sources is necessary; fidelity to the spread of opinion is a separate test, and no current review process applies it.

If that test becomes routine, the editor's role extends from checking facts to checking convergence. Editors are not currently trained to do that, nor is it part of their mandate. The system therefore has to make convergence visible to them, through a log they can read and question, rather than leave it to individual judgement under deadline. Regulators already warn that reliance on a small number of AI models and providers could lead market participants to take correlated positions (*Financial Stability Board, 2024*; *Bank of England, 2025*).

The same logic applies to interpretation: institutions could monitor dependence on a few readings of events, not only on a few models. Accumulated across many briefings and organisations, consensus-and-dissent logs would become an early-warning signal: a record of when commentary narrowed, and whether the minority view later proved right. In that sense, preserving dissent is a way of preserving the market's risk instinct.

Why dissent is worth keeping

Some would say that convergence is often simply correct. When the data is clear, independent economists reach the same view, and most contrarians are wrong; a system that elevates dissent risks amplifying noise and lending fringe views false weight.

That objection is partly right, and the design accepts it. London@11 does not judge which interpretation is correct, and it does not treat agreement as suspect. In the bench tests, it reported genuine consensus without faking balance and refused to elevate anonymous or unsupported counter-views; its credible-dissent rule admits only named analysts with stated reasoning. The independence test does a narrower job than separating right from wrong: it separates real agreement among separately reasoned sources

from apparent agreement produced by shared wire copy or a single underlying note, as when five outlets reduced to two sources.

The case for preserving dissent therefore does not rest on dissenters being right. It rests on two points. First, whether a consensus was correct can only be known later, and a summary that drops the minority view erases the evidence needed to judge it. Of the 2008 financial crisis, the US Financial Crisis Inquiry Commission found that "there were warning signs. The tragedy was that they were ignored or discounted" (*Financial Crisis Inquiry Commission, 2011*).

Second, keeping a dissenting view is cheap, but losing one can be costly: recording it takes a sentence and a log entry, while dropping one that proves right can be expensive. AI-driven narrowing is already measurable in science (*Hao et al., 2026*), and regulators warn of correlated behaviour in finance (*Bank of England, 2025*). A convergence that leaves no trace would make such losses harder to detect.

Over time, the record allows the question to be tested rather than argued: how far a dissenting view travelled before it was overridden, and how often it proved consequential. The same approach could apply wherever AI summarises expert opinion, not just in finance.

Conclusion

This article has argued that AI synthesis of financial commentary can make independent experts appear to agree, and that the safeguard must be built into the system rather than requested of it. London@11 indicates that this is practical: its source-independence test, credible-dissent rule and structured log held in live and controlled tests, though at small scale and under its designer's own evaluation.

The most important implication is that accuracy is not enough. Any tool that summarises expert opinion should be able to show which views it combined, which it set aside and why.

Three steps follow. First, a blind comparison, scored by independent editors, of London@11 against a standard AI summary of the same sources. Second, a controlled comparison of conventional and AI-assisted newsroom workflows on the same events, to measure whether AI narrows interpretation. Third, sustained operation: scaled over a year or two of continuous running, the accumulated record of flagged convergence and dissent could become a research instrument in its own right, a signal of where financial commentary reliably under-weighted minority views that later proved consequential, and where it didn't.

Developing a quantitative measure for this, scoring how far a dissenting view travelled before being overridden by consensus, is a natural further step, and one I am exploring separately as part of independent research into epistemic concentration risk in financial markets.

References

Bank of England (2025) *Financial Stability in Focus: Artificial intelligence in the financial system*. London: Bank of England, April. Available at: https://www.bankofengland.co.uk/-/media/boe/files/financial-stability-in-focus/2025/financial-stability-in-focus-artificial-intelligence-in-the-financial-system

Financial Crisis Inquiry Commission (2011) *The Financial Crisis Inquiry Report: Final Report of the National Commission on the Causes of the Financial and Economic Crisis in the United States*. Washington, DC: US Government Printing Office, January. Conclusions available at: https://cybercemetery.unt.edu/archive/fcic/20110310173539/https://c0182732.cdn1.cloudfiles.rackspacecloud.com/fcic_final_report_conclusions.pdf

Financial Stability Board (2024) *The Financial Stability Implications of Artificial Intelligence*. Basel: FSB, 14 November. Available at: https://www.fsb.org/uploads/P14112024.pdf

Hao, Q., Xu, F., Li, Y. and Evans, J. (2026) 'Artificial intelligence tools expand scientists' impact but contract science's focus', *Nature*, 649, pp. 1237–1243. Available at: https://doi.org/10.1038/s41586-025-09922-y

Imas, A. and Manyika, J. (2026) 'Bending the Curve of Discovery: AI in Science Today and Tomorrow', *DeepMind Institute*, 8 October. Available at: https://institute.deepmind.com/essays/bending-the-curve-of-discovery-ai-in-science-today-and-tomorrow/

Use of AI

London@11 runs on Anthropic's Claude; the early tests described in "Evidence of convergence" used ChatGPT Enterprise. Claude also reviewed live drafts from 29 September 2026, proposing policy changes I accepted or rejected, and built, ran and scored the bench tests against criteria set in advance. I used Claude to refine and review the article content and check facts and sources. I am accountable for the final text.